Task 23

Disaster Recovery Setup
Enable:

- S3 Cross-Region Replication
- RDS Read Replica
- AMI backup
- Simulate region failure

Create **destination bucket** in another region

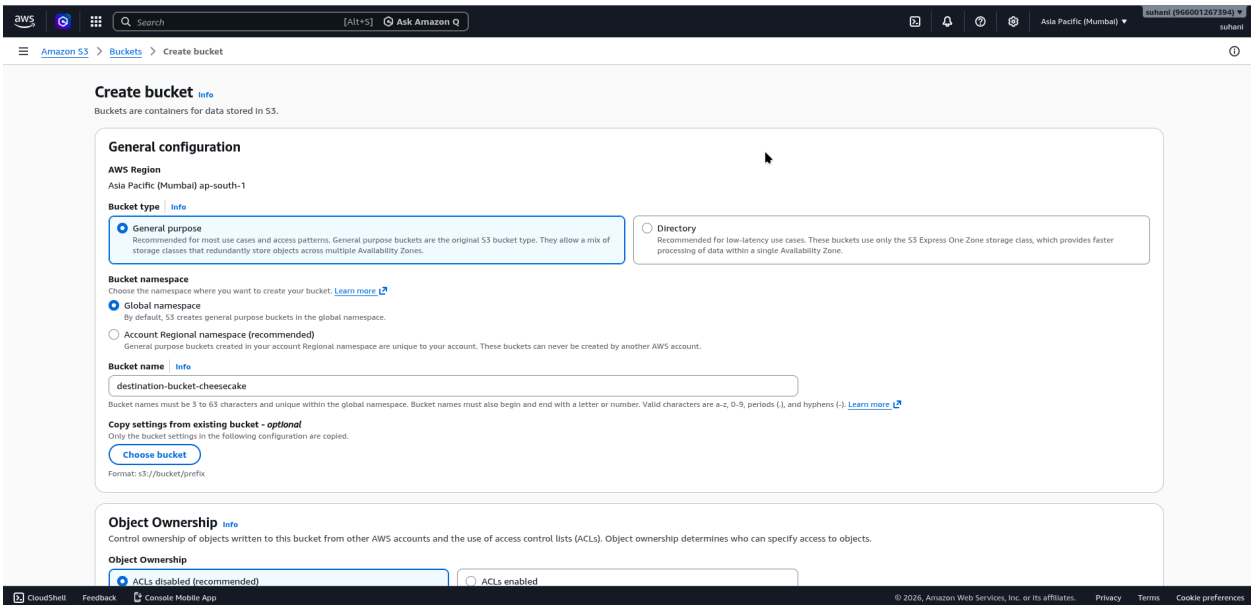

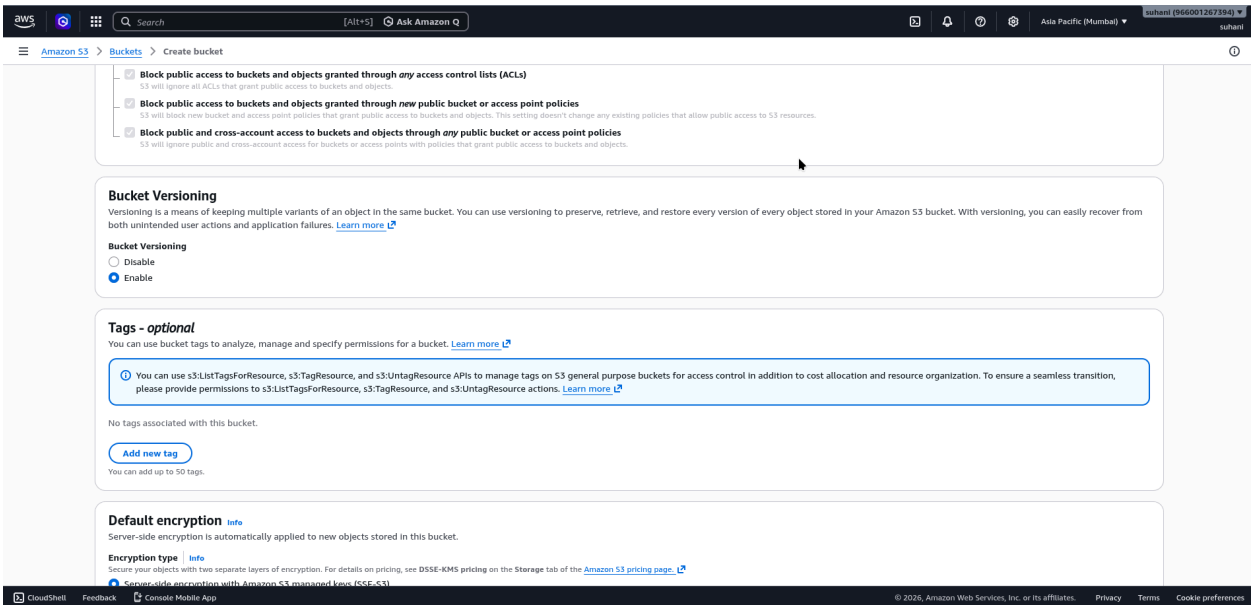

Source bucket (your main):

- **Properties → Versioning → Enable**

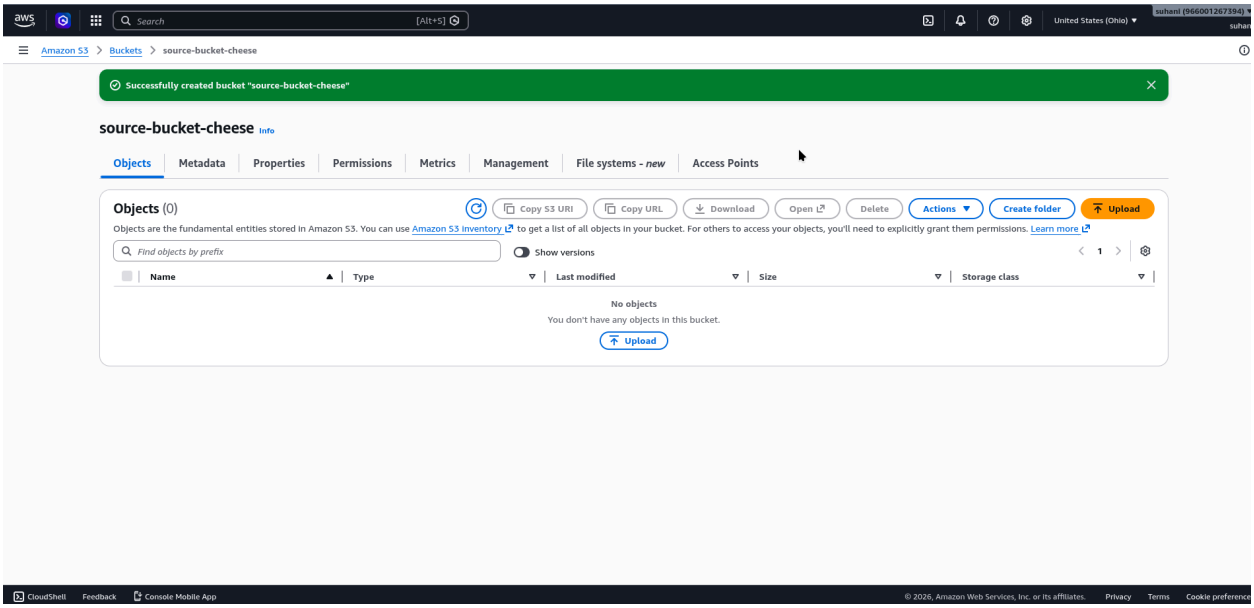

1. Source bucket → **Management → Replication rules → Create rule**
    - Rule name: `crr-rule`
    - Scope: **All objects** (or prefix)
    - Destination:

- Bucket: `my-dr-bucket-ap-southeast-1`

- Create/choose IAM role (let AWS create)

👉 Save

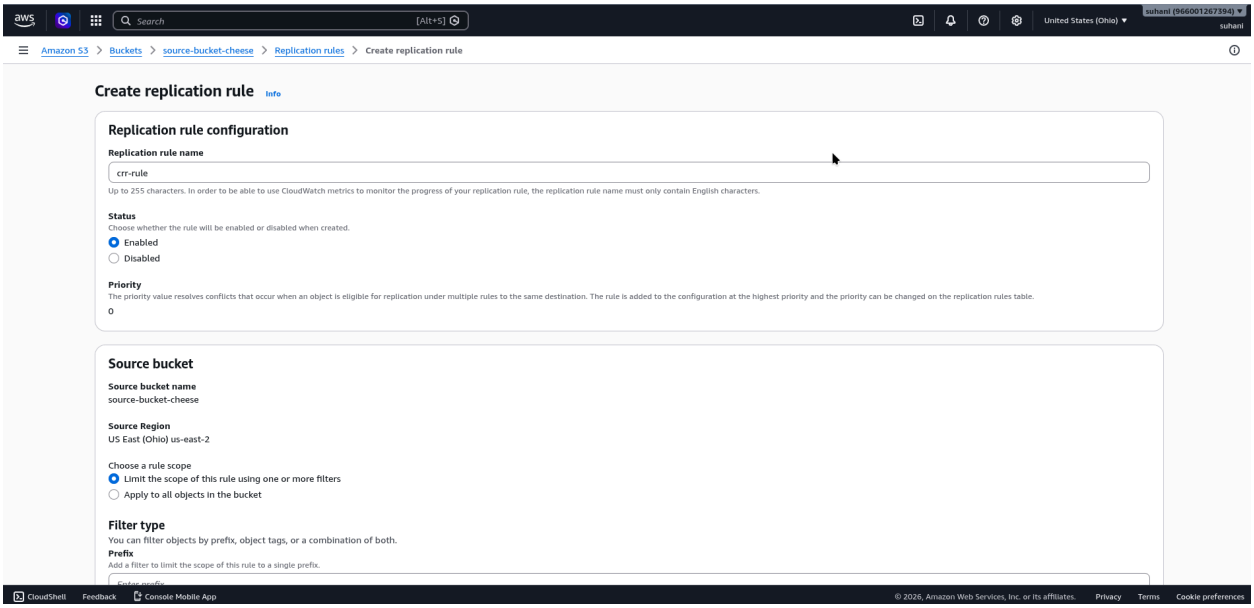

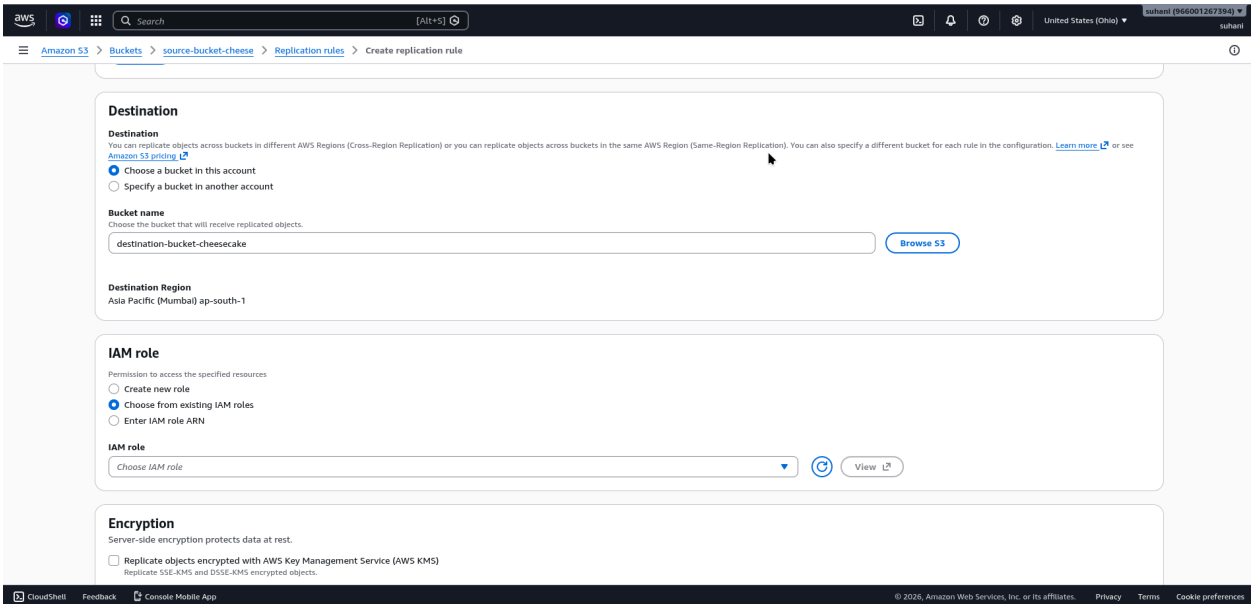

IAM ROLE AND PREFIX IMP

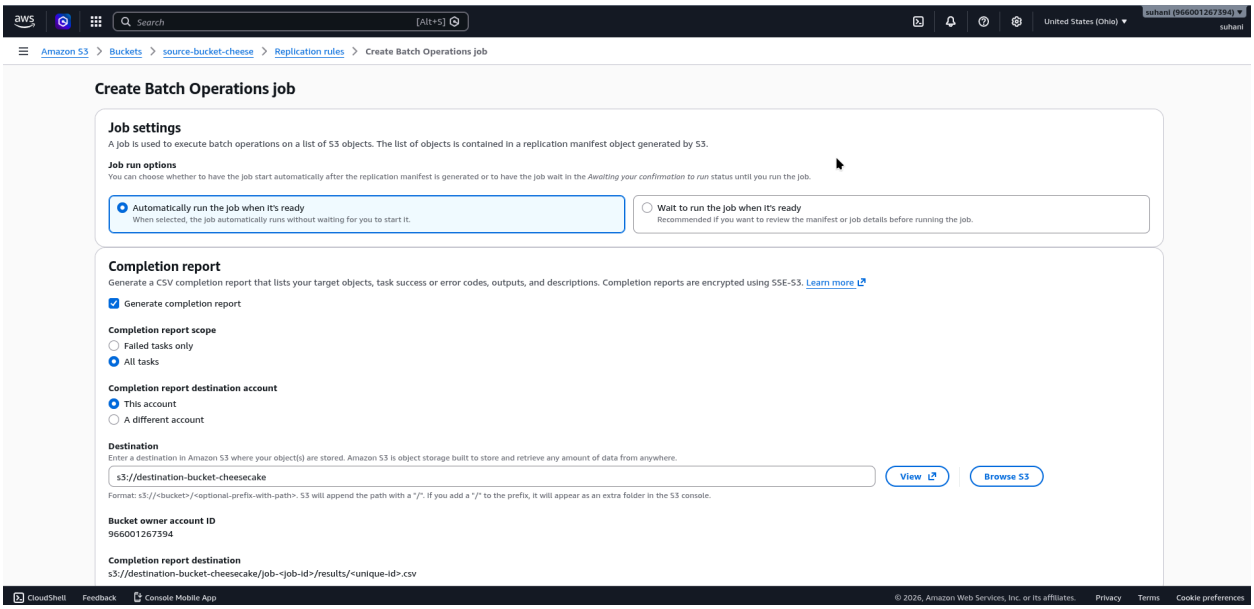
Search
[Alt+S]
United States (Ohio)
suhani (966001267394)
suhani
Amazon S3 > Buckets > source-bucket-cheese > Replication rules > Create Batch Operations job
Create Batch Operations job
Job settings
A Job is used to execute batch operations on a list of S3 objects. The list of objects is contained in a replication manifest object generated by S3.
Job run options
You can choose whether to have the Job start automatically after the replication manifest is generated or to have the Job wait in the Awaiting your confirmation to run status until you run the Job.
Automatically run the job when it's ready
When selected, the job automatically runs without waiting for you to start it.
Wait to run the job when it's ready
Recommended if you want to review the manifest or job details before running the job.
Completion report
Generate a CSV completion report that lists your target objects, task success or error codes, outputs, and descriptions. Completion reports are encrypted using SSE-S3. Learn more
Generate completion report
Completion report scope
Failed tasks only
All tasks
Completion report destination account
This account
A different account
Destination
Enter a destination in Amazon S3 where your object(s) are stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.
s3://destination-bucket-cheesecake
View
Browse S3
Format: s3://<bucket>/<optional-prefix-with-path>. S3 will append the path with a "/". If you add a "/" to the prefix, it will appear as an extra folder in the S3 console.
Bucket owner account ID
966001267394
Completion report destination
s3://destination-bucket-cheesecake/job-<job-id>/results/<unique-id>.csv
CloudShell
Feedback
Console Mobile App
© 2026, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

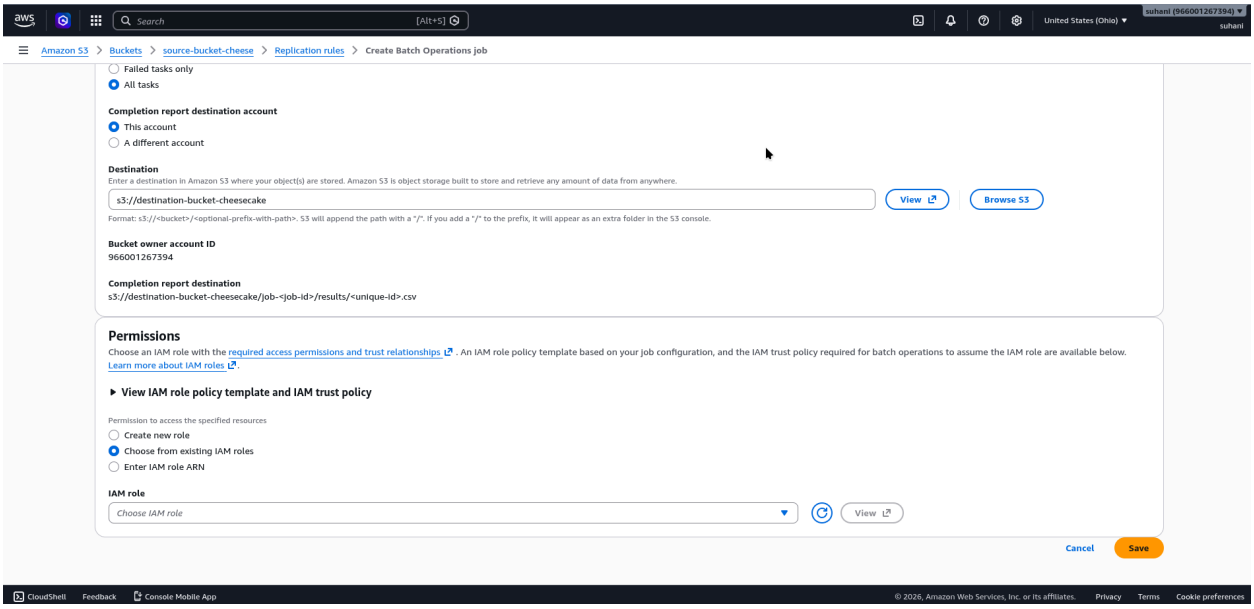
Search
[Alt+S]
United States (Ohio)
suhani (966001267394)
suhani
Amazon S3 > Buckets > source-bucket-cheese > Replication rules > Create Batch Operations job
Failed tasks only
All tasks
Completion report destination account
This account
A different account
Destination
Enter a destination in Amazon S3 where your object(s) are stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.
s3://destination-bucket-cheesecake
View
Browse S3
Format: s3://<bucket>/<optional-prefix-with-path>. S3 will append the path with a "/". If you add a "/" to the prefix, it will appear as an extra folder in the S3 console.
Bucket owner account ID
966001267394
Completion report destination
s3://destination-bucket-cheesecake/job-<job-id>/results/<unique-id>.csv
Permissions
Choose an IAM role with the required access permissions and trust relationships . An IAM role policy template based on your job configuration, and the IAM trust policy required for batch operations to assume the IAM role are available below.
Learn more about IAM roles.
View IAM role policy template and IAM trust policy
Permission to access the specified resources
Create new role
Choose from existing IAM roles
Enter IAM role ARN
IAM role
Choose IAM role
View
Cancel
Save
CloudShell
Feedback
Console Mobile App
© 2026, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

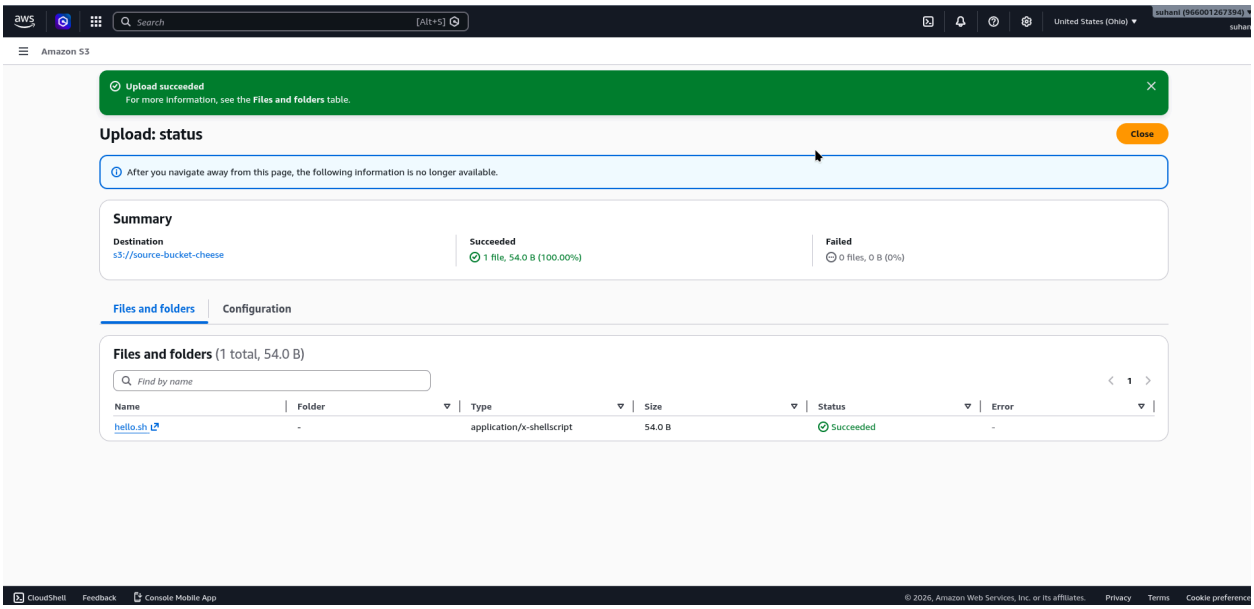

uploaded in source

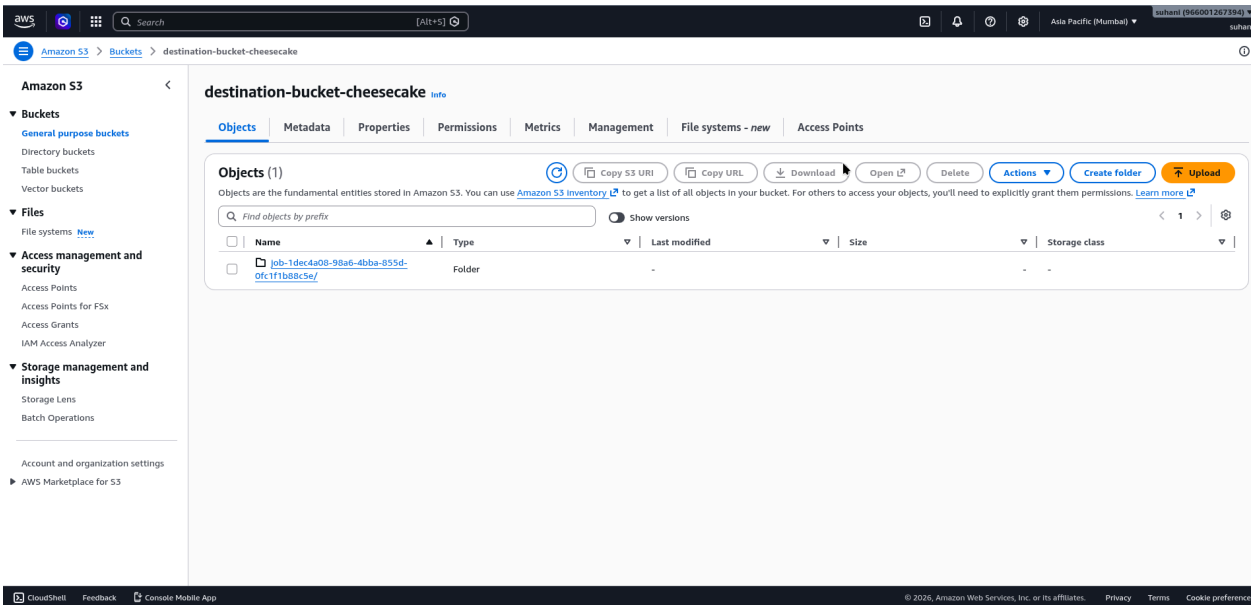

found as object in destination

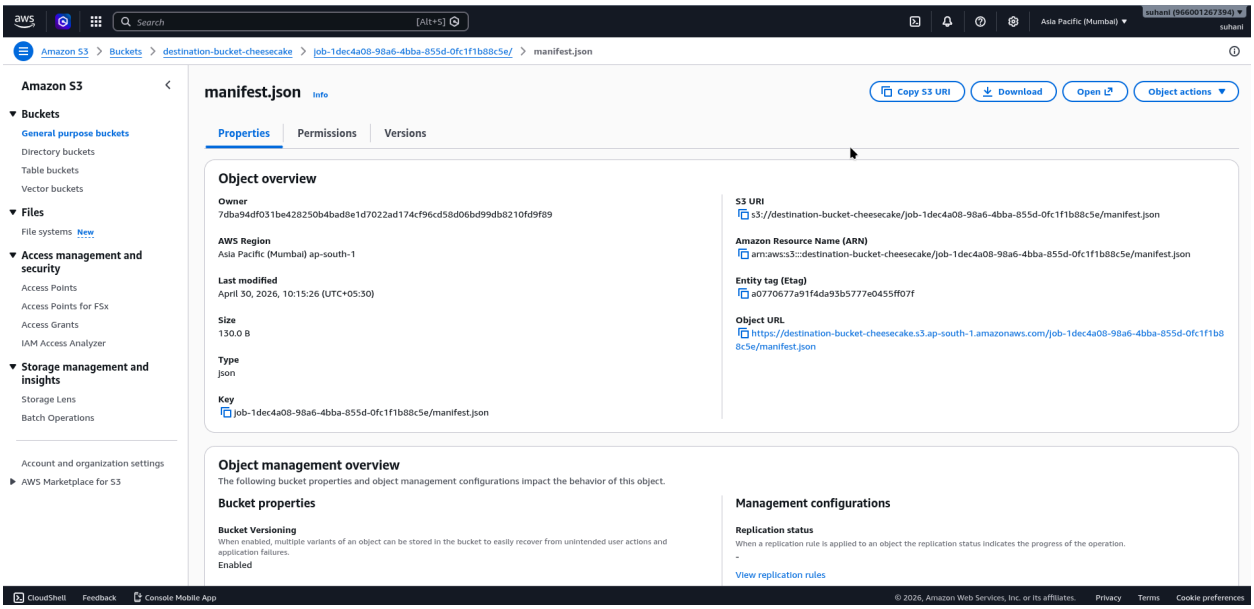

2) RDS Read Replica

Steps

1. Go to
   👉 Amazon RDS Console
2. Select your primary DB → **Actions** → **Create read replica**

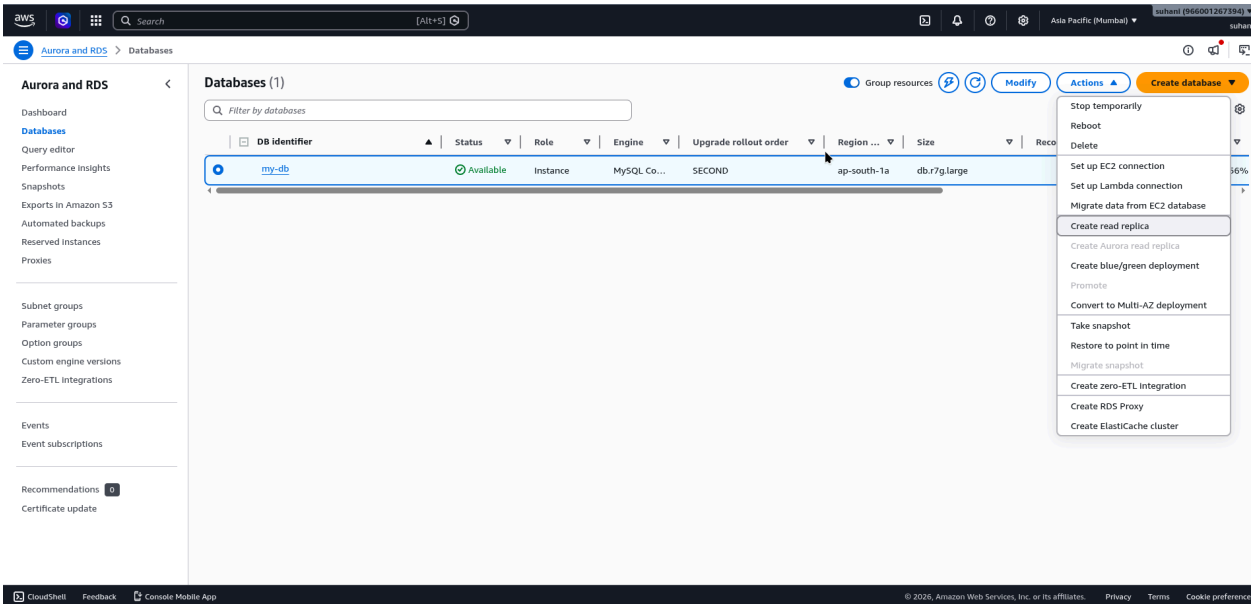

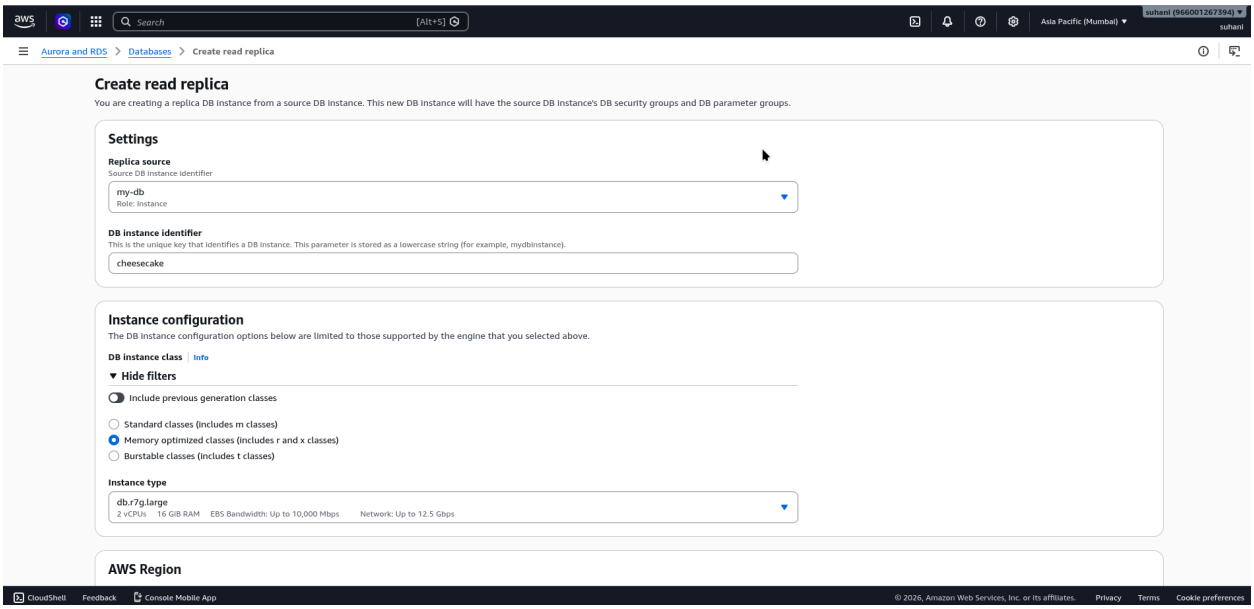
Search
[Alt+S]
Asia Pacific (Mumbai)
suhani (966001267394)
suhani
Aurora and RDS > Databases > Create read replica
Create read replica
You are creating a replica DB instance from a source DB instance. This new DB instance will have the source DB instance's DB security groups and DB parameter groups.
Settings
Replica source
Source DB instance identifier
my-db
Role: Instance
DB instance identifier
This is the unique key that identifies a DB instance. This parameter is stored as a lowercase string (for example, mydbinstance).
cheesecake
Instance configuration
The DB instance configuration options below are limited to those supported by the engine that you selected above.
DB instance class | Info
Hide filters
Include previous generation classes
Standard classes (includes m classes)
Memory optimized classes (includes r and x classes)
Burstable classes (includes t classes)
Instance type
db.r7g.large
2 vCPUs 16 GiB RAM EBS Bandwidth: Up to 10,000 Mbps Network: Up to 12.5 Gbps
AWS Region
CloudShell
Feedback
Console Mobile App
© 2026, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

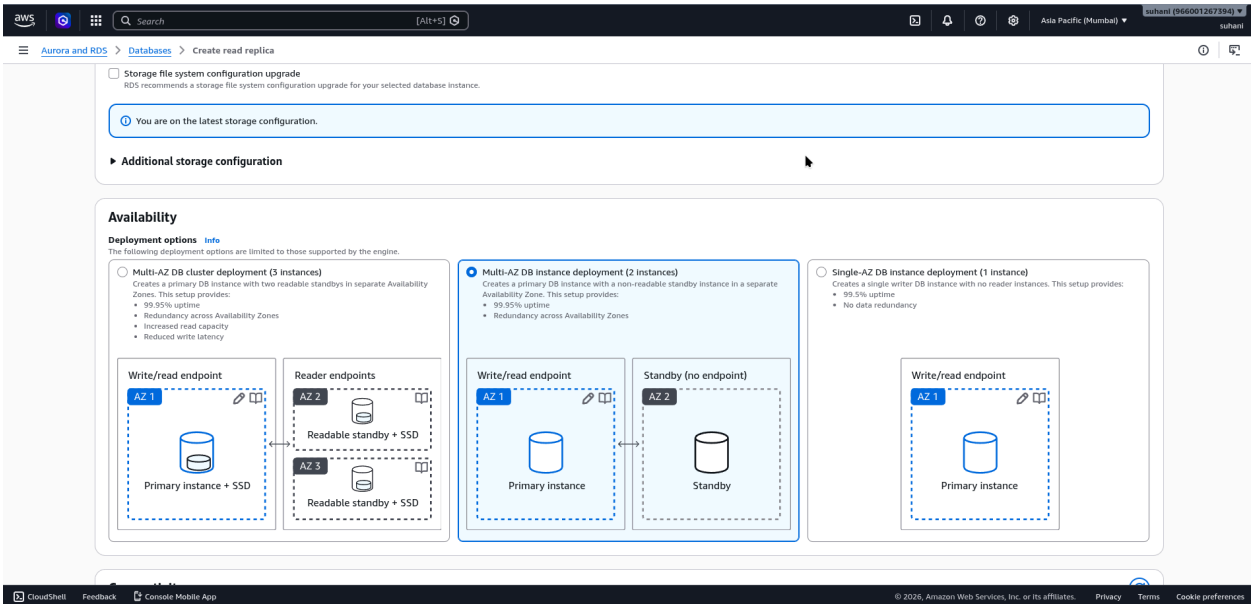
Search
[Alt+S]
Asia Pacific (Mumbai)
suhani (966001267394)
suhani
Aurora and RDS > Databases > Create read replica
Storage file system configuration upgrade
RDS recommends a storage file system configuration upgrade for your selected database instance.
You are on the latest storage configuration.
Additional storage configuration
Availability
Deployment options Info
The following deployment options are limited to those supported by the engine.
Multi-AZ DB cluster deployment (3 instances)
Creates a primary DB instance with two readable standbys in separate Availability Zones. This setup provides:
99.95% uptime
Redundancy across Availability Zones
Increased read capacity
Reduced write latency
Write/read endpoint
AZ 1
Primary instance + SSD
Reader endpoints
AZ 2
Readable standby + SSD
AZ 3
Readable standby + SSD
Multi-AZ DB instance deployment (2 instances)
Creates a primary DB instance with a non-readable standby instance in a separate Availability Zone. This setup provides:
99.95% uptime
Redundancy across Availability Zones
Write/read endpoint
AZ 1
Primary instance
Standby (no endpoint)
AZ 2
Standby
Single-AZ DB instance deployment (1 instance)
Creates a single writer DB instance with no reader instances. This setup provides:
99.5% uptime
No data redundancy
Write/read endpoint
AZ 1
Primary instance
CloudShell
Feedback
Console Mobile App
© 2026, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

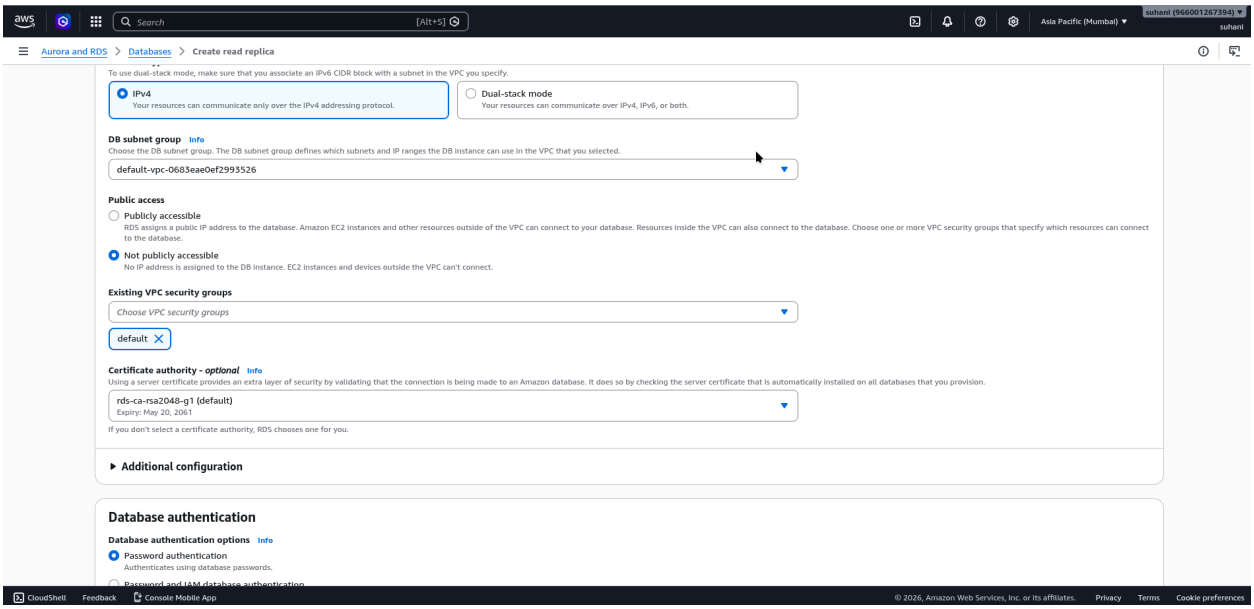

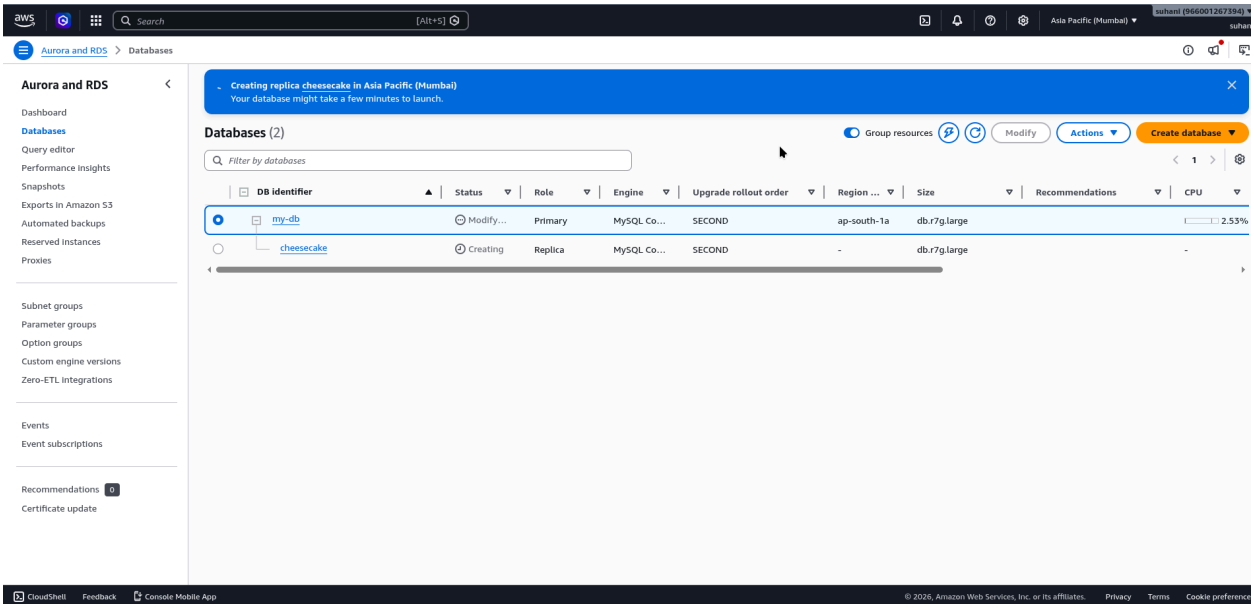

- Status becomes **Available**
- Shows as **Replica of primary**

AMI Backup (EC2)

Steps

Go to

👉 Amazon EC2 Console

1. Select your EC2 instance

Fill:

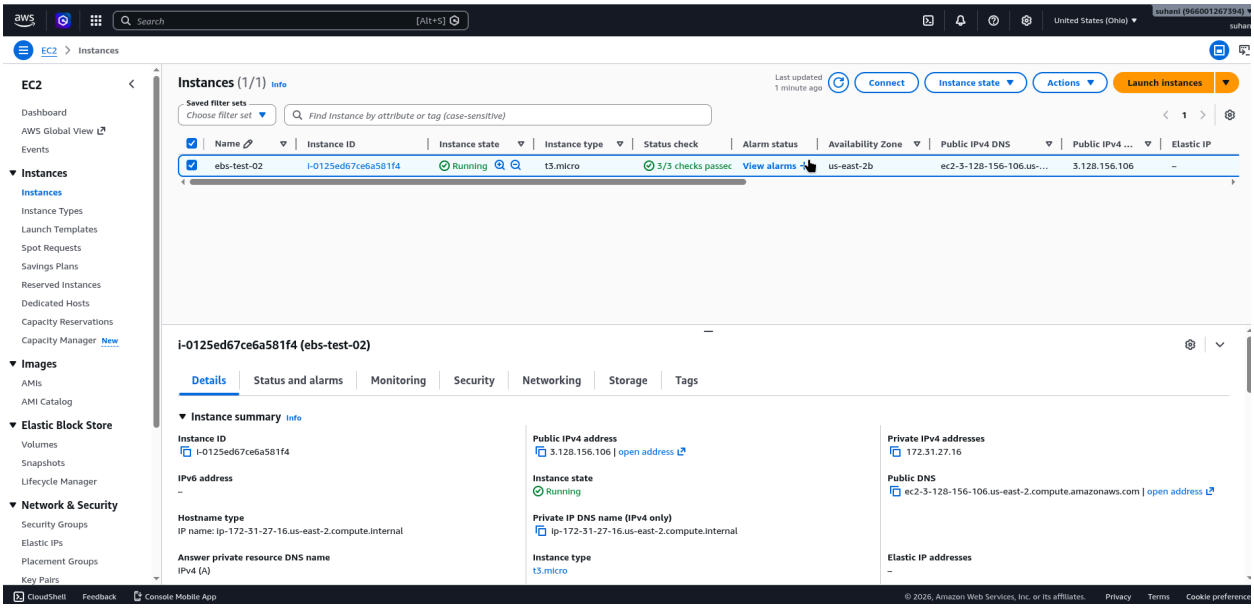

- Name: `my-app-backup`

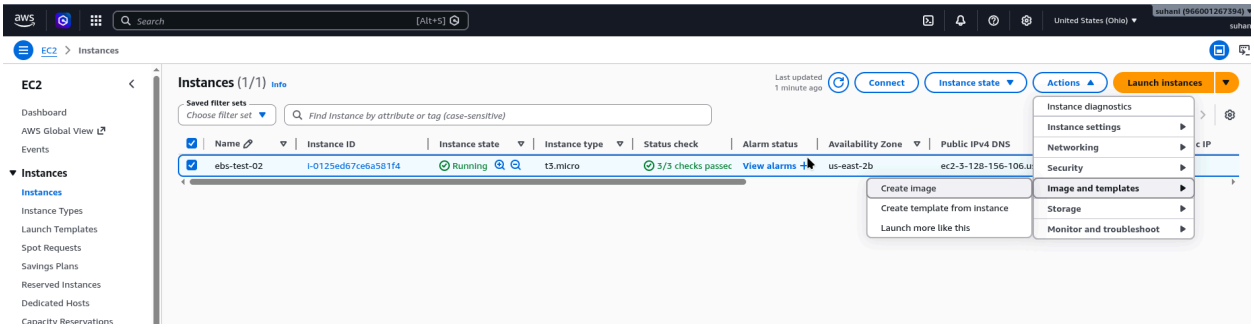

Validation

- Go to **AMIs** → **Owned by me**
- Status → **Available**

4) **Simulate Region Failure**

S3

- Delete object from source
- Show it exists in DR bucket

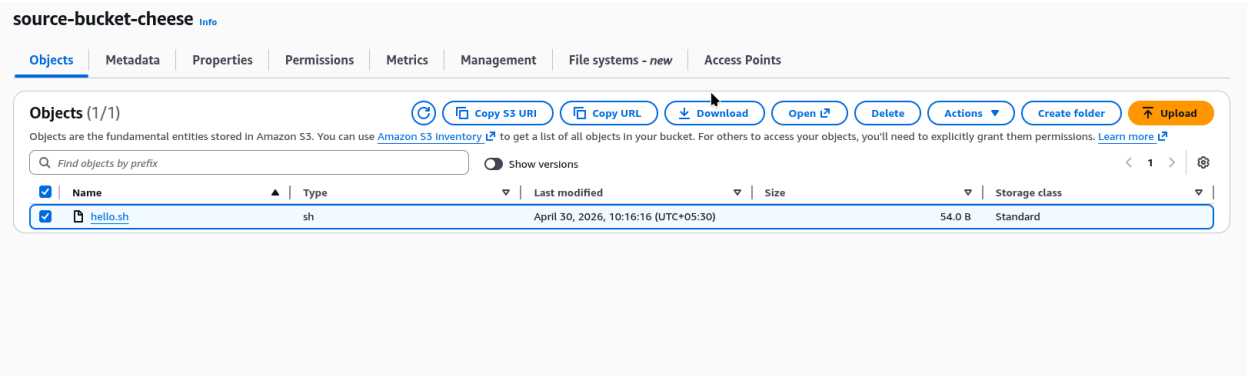

i will delete this shell script

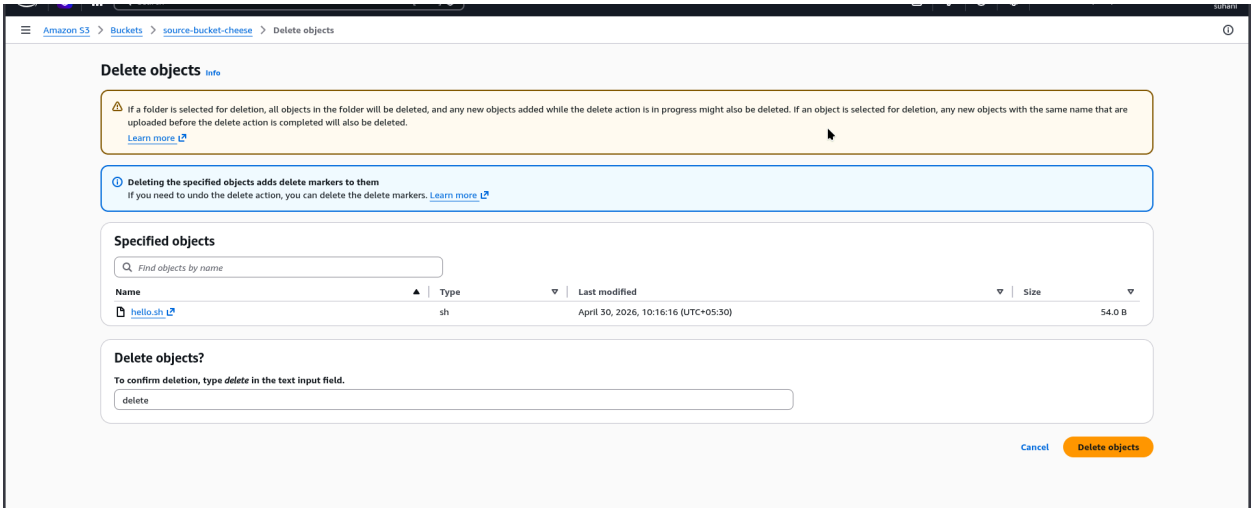

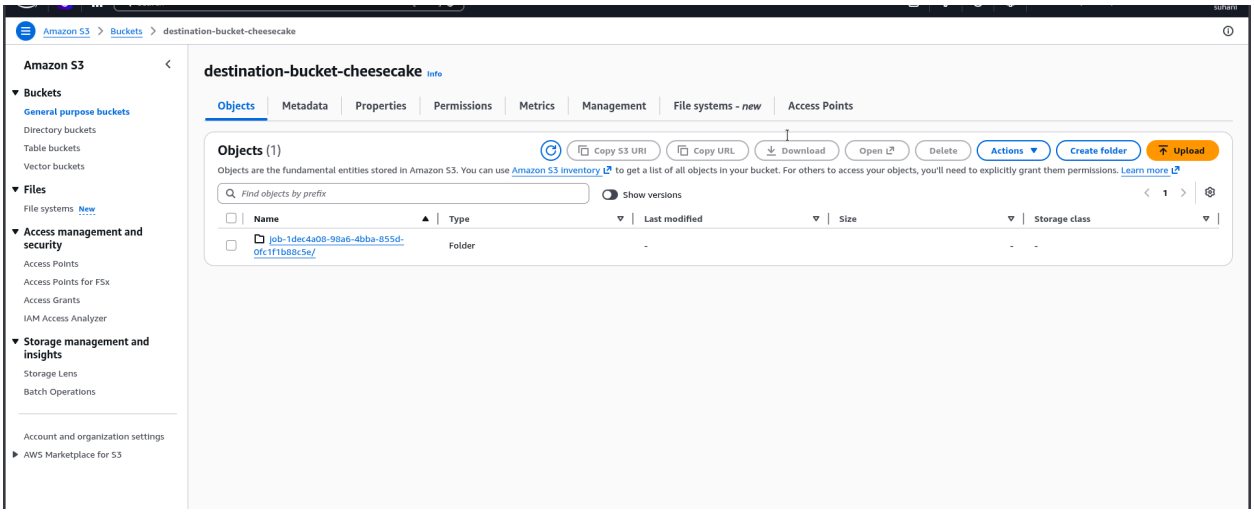

but the object is still shown in destination